

# Reinforcement Learning

Jens Kober

[j.kober@tudelft.nl](mailto:j.kober@tudelft.nl)

# Outline

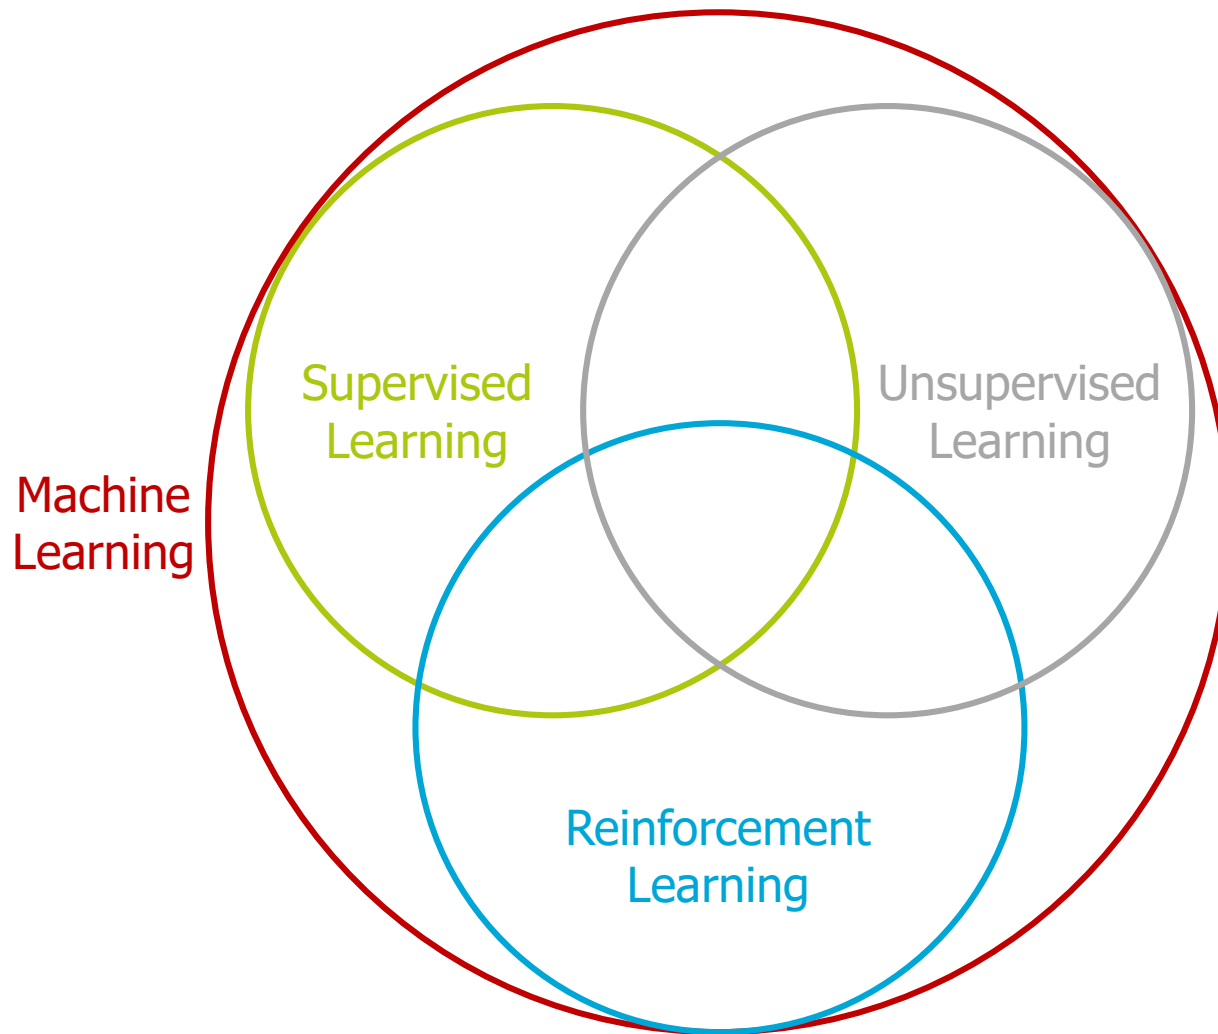
- Introduction
- Notation & basic concepts
- Types of RL algorithms
  - Optimal control
  - Value function methods
  - Policy search
- Summary & outlook

# Types of Machine Learning

- Supervised learning
- Unsupervised learning
- Semi-supervised learning
- Reinforcement learning

- Definitions?
- Differences?
- Examples?

# Types of Machine Learning



# Reinforcement Learning

- Learning by trial & error



- Learning by rewards & punishments



# Outline

- Introduction
- Notation & basic concepts
- Types of RL algorithms
  - Optimal control
  - Value function methods
  - Policy search
- Summary & outlook


# Classical RL Example

- Maze



# Classical RL Example


- Maze = grid world







# Classical RL Example

- Strategy = policy






# Classical RL Example

- Deterministic grid world

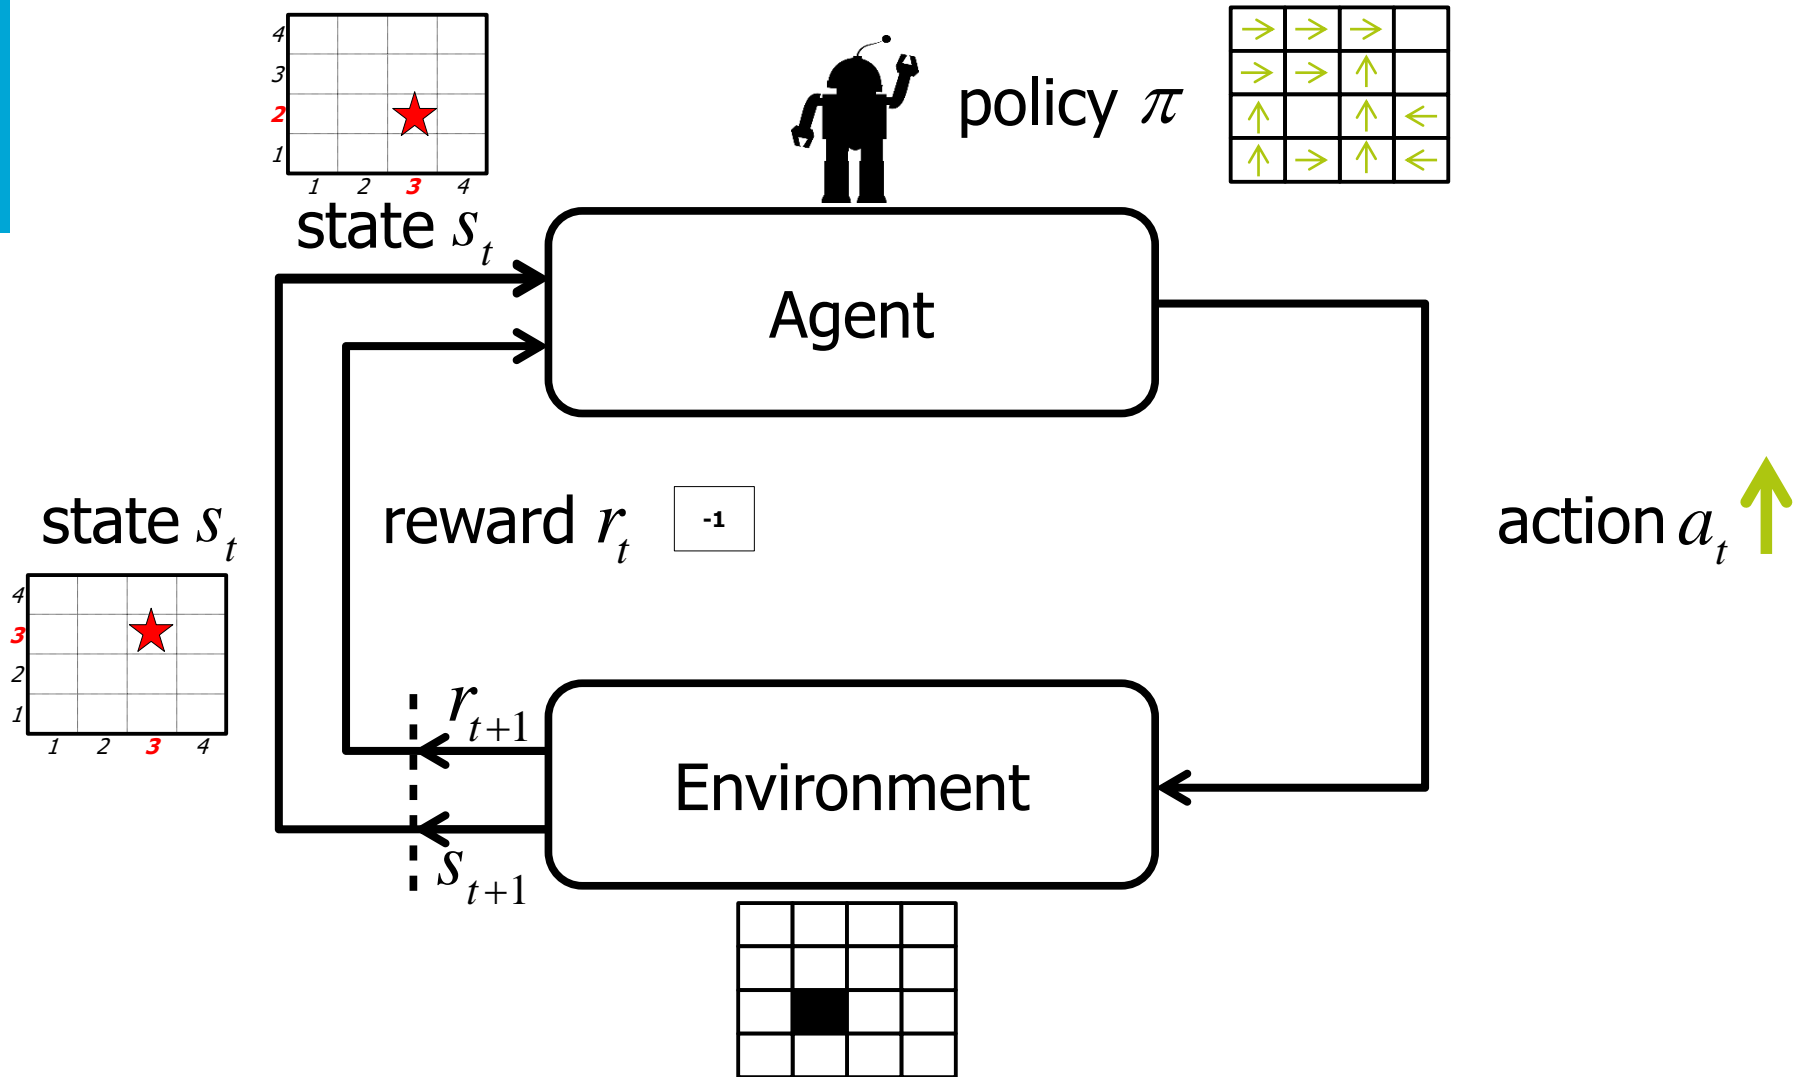


transition probabilities


- Stochastic grid world



# Formal Notation



# Formal Notation



- $s \in \mathcal{S}$  set of states (discrete or continuous)
- $a \in \mathcal{A}$  set of actions (discrete or continuous)
- $r$  reward  $\mathcal{R}_{ss'}^a = E \{ r_{t+1} \mid a_t = a, s_t = s, s_{t+1} = s' \}$
- $\pi(s) = a$  policy

Goal: find  $\pi^*$  maximizing return  $R$

# Return $R$

- Finite Horizon  $H$

$$R_t = r_{t+1} + r_{t+2} + \dots + r_{t+H} = \sum_{h=1}^H r_{t+h}$$

- Discounted  $0 \leq \gamma < 1$

$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+2} + \dots = \sum_{h=0}^{\infty} \gamma^h r_{t+h+1}$$



- Average

$$R_t = \lim_{H \rightarrow \infty} \left( \frac{1}{H} \sum_{h=1}^H r_{t+h} \right)$$

# Markov Property

- Transition probability

$$P(s_{t+1} | a_t, s_t, \cancel{a_{t-1}, s_{t-1}, \dots, a_1, s_1})$$

## ➤ Markov property

$$\mathcal{P}_{ss'}^a = P(s_{t+1} = s' | s_t = s, a_t = a)$$

- Markov decision process (MDP)

$$\mathcal{S}, \mathcal{A}, \mathcal{P}_{ss'}^a, \mathcal{R}_{ss'}^a, \gamma$$

# Bellman Principle of Optimality

“An optimal policy has the property that whatever the initial state and initial decision are, the remaining decisions must constitute an optimal policy with regard to the state resulting from the first decision.”

Bellman, 1957

- Dynamic Programming
- Optimal Control

# Value Functions

- State value function

$$V^{\pi}(s) = E^{\pi} \left\{ \sum_{h=0}^{\infty} \gamma^h r_{t+h+1} \mid s_t = s \right\} = \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^{\pi} \left[ \mathcal{R}_{ss'}^{\pi} + \gamma V^{\pi}(s') \right]$$

$$V^*(s) = \max_{a \in \mathcal{A}} \left( \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma V^*(s') \right] \right)$$

- State-action value function

$$Q^{\pi}(s, a) = \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma V^{\pi}(s') \right]$$

$$\begin{aligned} Q^*(s, a) &= \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma V^*(s') \right] \\ &= \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma \left( \max_{a' \in \mathcal{A}} Q^*(s', a') \right) \right] \end{aligned}$$



# Optimal Policy

- State-action value function

$$Q^*(s, a) = \sum_{s' \in S} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma \left( \max_{a' \in \mathcal{A}} Q^*(s', a') \right) \right]$$

➤  $\pi^*(s) = \arg \max_{a \in \mathcal{A}} (Q^*(s, a))$

- State value function

$$V^*(s) = \max_{a \in \mathcal{A}} \left( \sum_{s' \in S} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma V^*(s') \right] \right)$$

➤  $\pi^*(s) = \arg \max_{a \in \mathcal{A}} \left( \sum_{s' \in S} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma V^*(s') \right] \right)$

# Optimal State-Action Value Function

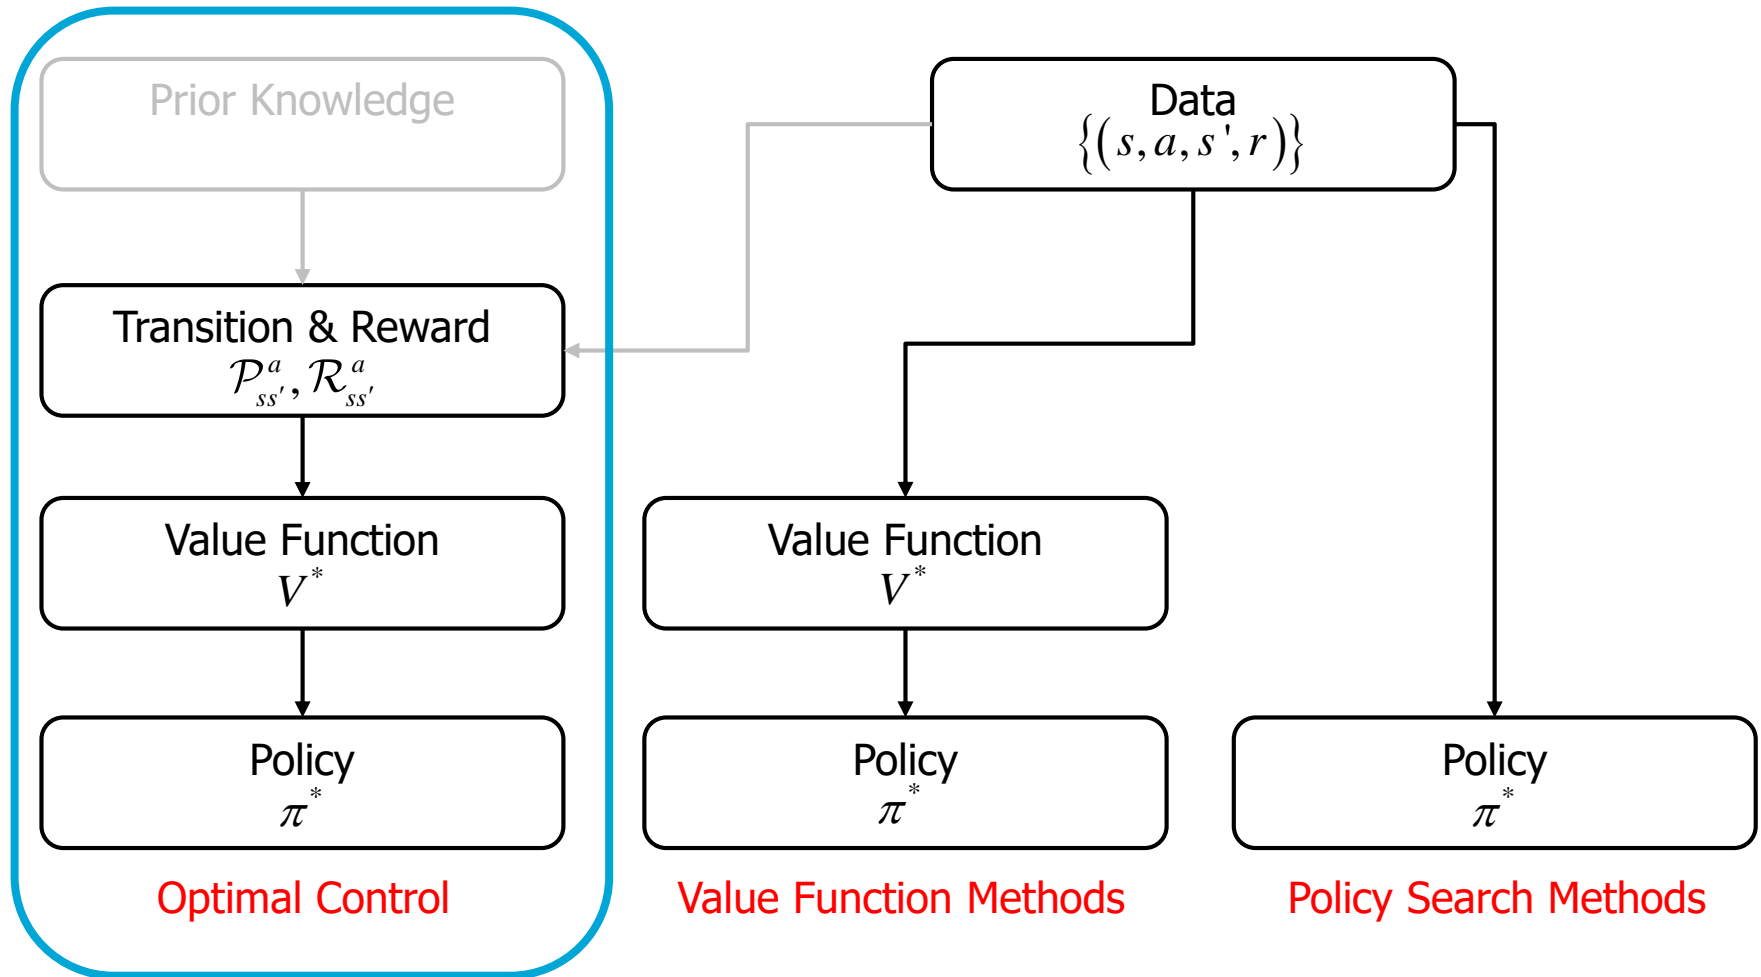
- max. 10 steps, stop at fields  $\pm 10$ , stay at walls, deterministic,  $\gamma = 1$ , reward: -1 per step,  $\pm 10$  for reaching states

		6	7	8	
4	5	-1 7	6 -1 8	7 -1 9	+10
		5	6	7	
3	5	6 -1 6	5 -1 7	6 -1 -11	-10
		4	6	6	
2	4	5 -1 4		6 -1 5	-11 6 -1 5
		3		5	4
1	3	4 -1 4	3 -1 5	4 -1 4	5 -1 4
		3	4	5	4
		1	2	3	4

# Outline

- Introduction
- Notation & basic concepts
- Types of RL algorithms
  - Optimal control
  - Value function methods
  - Policy search
- Summary & outlook

# Types of RL Algorithms



# Optimal Control

- **Model-based:** requires models of reward and transition functions
- Fast (no real world interactions required)
- Models can be also be learned
- RL often takes advantage of model errors
- Examples:
  - Policy Iteration
  - Value Iteration

# Value Iteration

- Bellman optimality equation

$$Q^*(s, a) = \sum_{s' \in S} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma \left( \max_{a' \in \mathcal{A}} Q^*(s', a') \right) \right]$$

- Turn into an **iterative update**

## Q-Iteration

**repeat** at each iteration  $i$

**for all**  $s, a$  **do**

$$Q_{i+1}(s, a) \leftarrow \sum_{s' \in S} \mathcal{P}_{ss'}^a \left[ \mathcal{R}_{ss'}^a + \gamma \left( \max_{a' \in \mathcal{A}} Q_i(s', a') \right) \right]$$

**end for**

**until** convergence to  $Q^*$

- Once  $Q^*$  available  $\pi^*(s) = \arg \max_{a \in \mathcal{A}} (Q^*(s, a))$

# 1<sup>st</sup> Iteration

$Q_0$

4	0	-1	0	0	-1	0	0	-1	0	+10
	0			0				0		
3	0	-1	0	0	-1	0	0	-1	0	-10
	0			0				0		
2	0	-1	0				0	-1	0	0
	0						0			
1	0	-1	0	0	-1	0	0	-1	0	0
	0			0			0			
	1		2		3		4			

$Q_1$

		-1	-1	-1				
4	-1	-1	-1	-1	-1	-1	9	+10
		-1		-1			-1	
	-1	-1	-1	-1	-1	-1	-1	-10
		-1		-1			-1	
3	-1	-1	-1	-1	-1	-1	-1	
		-1		-1			-1	
	-1	-1	-1			-1	-1	-1
		-1		-1			-1	
2	-1	-1	-1			-1	-1	-1
		-1		-1			-1	
	-1	-1	-1	-1	-1	-1	-1	-1
		-1		-1			-1	
1	-1	-1	-1	-1	-1	-1	-1	-1
		-1		-1			-1	
	1		2		3		4	

# 2<sup>nd</sup> Iteration

$Q_1$

4	-1 -1 -1	-1 -1 -1	-1 -1 9	<b>+10</b>
3	-1 -1 -1	-1 -1 -1	-1 -1 -11	<b>-10</b>
2	-1 -1 -1		-1 -1 -1	-11 -1 -1
1	-1 -1 -1	-1 -1 -1	-1 -1 -1	-1 -1 -1
	1	2	3	4

$Q_2$

4	-2 -2 -2	-2 -2 8	8 -1 -1 9	<b>+10</b>
3	-2 -2 -2	-2 -2 -2	-2 -1 -11	<b>-10</b>
2	-2 -2 -2		-2 -1 -2	-11 -2 -2
1	-2 -2 -2	-2 -2 -2	-2 -1 -2	-2 -2 -2
	1	2	3	4



# 3<sup>rd</sup> Iteration

$Q_2$

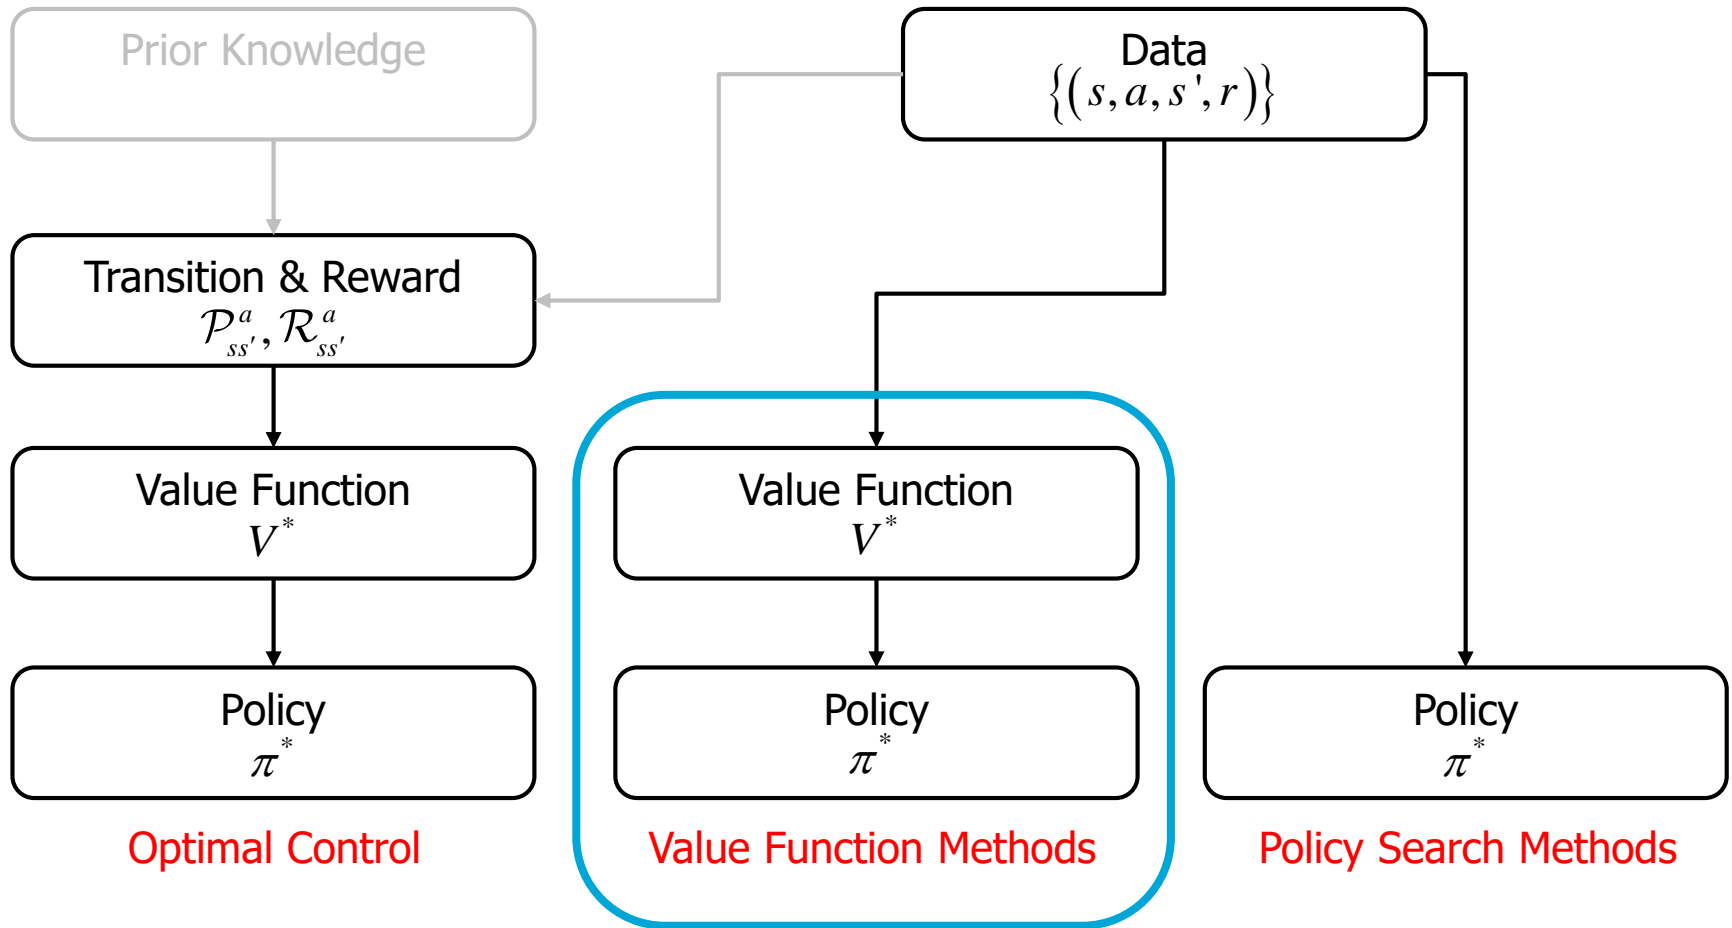
4	-2 -2 -2	-2 -1 -2	8 8 -1	<b>+10</b>
3	-2 -2 -2	-2 -1 -2	8 -1 -11	<b>-10</b>
2	-2 -2 -2		-2 -1 -2	-11 -2 -2
1	-2 -2 -2	-2 -1 -2	-2 -1 -2	-2 -1 -2
	1	2	3	4

$Q_3$

4	-3 -3 -3	7 -1 8	8 -1 9	<b>+10</b>
3	-3 -3 -3	7 -1 7	8 -1 -11	<b>-10</b>
2	-3 -3 -3		7 -1 -3	-11 -3 -3
1	-3 -3 -3	-3 -1 -3	-3 -1 -3	-3 -1 -3
	1	2	3	4

<https://youtu.be/VCdxqn0fcnE>

# Types of RL Algorithms



# Value Function Methods

- **Model-free:** (implicitly) learns models of reward and transition functions
- Global optimum (if everything is explored)
- Policy has global dependence
- Examples:
  - Monte Carlo
  - TD( $\lambda$ )
  - Q-learning
  - SARSA

# Temporal Difference Learning

- Use sample based estimate to update the value function

$$Q(s, a) \leftarrow (1 - \alpha) Q(s, a) + \alpha \left[ r + \gamma \max_{a' \in \mathcal{A}} Q(s', a') \right]$$

learning rate

- Equivalently

$$Q(s, a) \leftarrow Q(s, a) + \alpha \underbrace{\left[ r + \gamma \max_{a' \in \mathcal{A}} Q(s', a') - Q(s, a) \right]}_{\text{temporal difference error}}$$

# Exploration

- Which action  $a$  to pick?
- Assume  $Q$  is optimal (greedy action)

$$a \leftarrow \arg \max_{a \in \mathcal{A}} (Q^*(s, a))$$

- But  $Q$  is not optimal (yet)!
- We need to explore

$\varepsilon$ -greedy policy:

- With probability  $\varepsilon$  pick a random action (exploration)  
 $a \leftarrow \text{rand}(\mathcal{A})$
- With probability  $(1 - \varepsilon)$  pick the greedy action (exploitation)

$$a \leftarrow \arg \max_{a \in \mathcal{A}} (Q^*(s, a))$$

# Q-Learning

## Q-Learning

### loop

observe state  $s$

$a \leftarrow \arg \max_{a \in \mathcal{A}} (Q(s, a))$

with probability  $\varepsilon$ ,  $a \leftarrow \text{rand}(\mathcal{A})$

apply  $a$ , observe  $r$  and  $s'$

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \max_{a' \in \mathcal{A}} Q(s', a') - Q(s, a) \right]$$

**until** convergence


# Properties of Discussed Methods

- Advantages
  - Generic framework
  - Very few assumptions
  - Guaranteed to converge to optimum
- Disadvantages
  - Can take lot of iterations
  - Infeasible for continuous states/actions




# Dealing with Continuous States/Actions

- Discretization



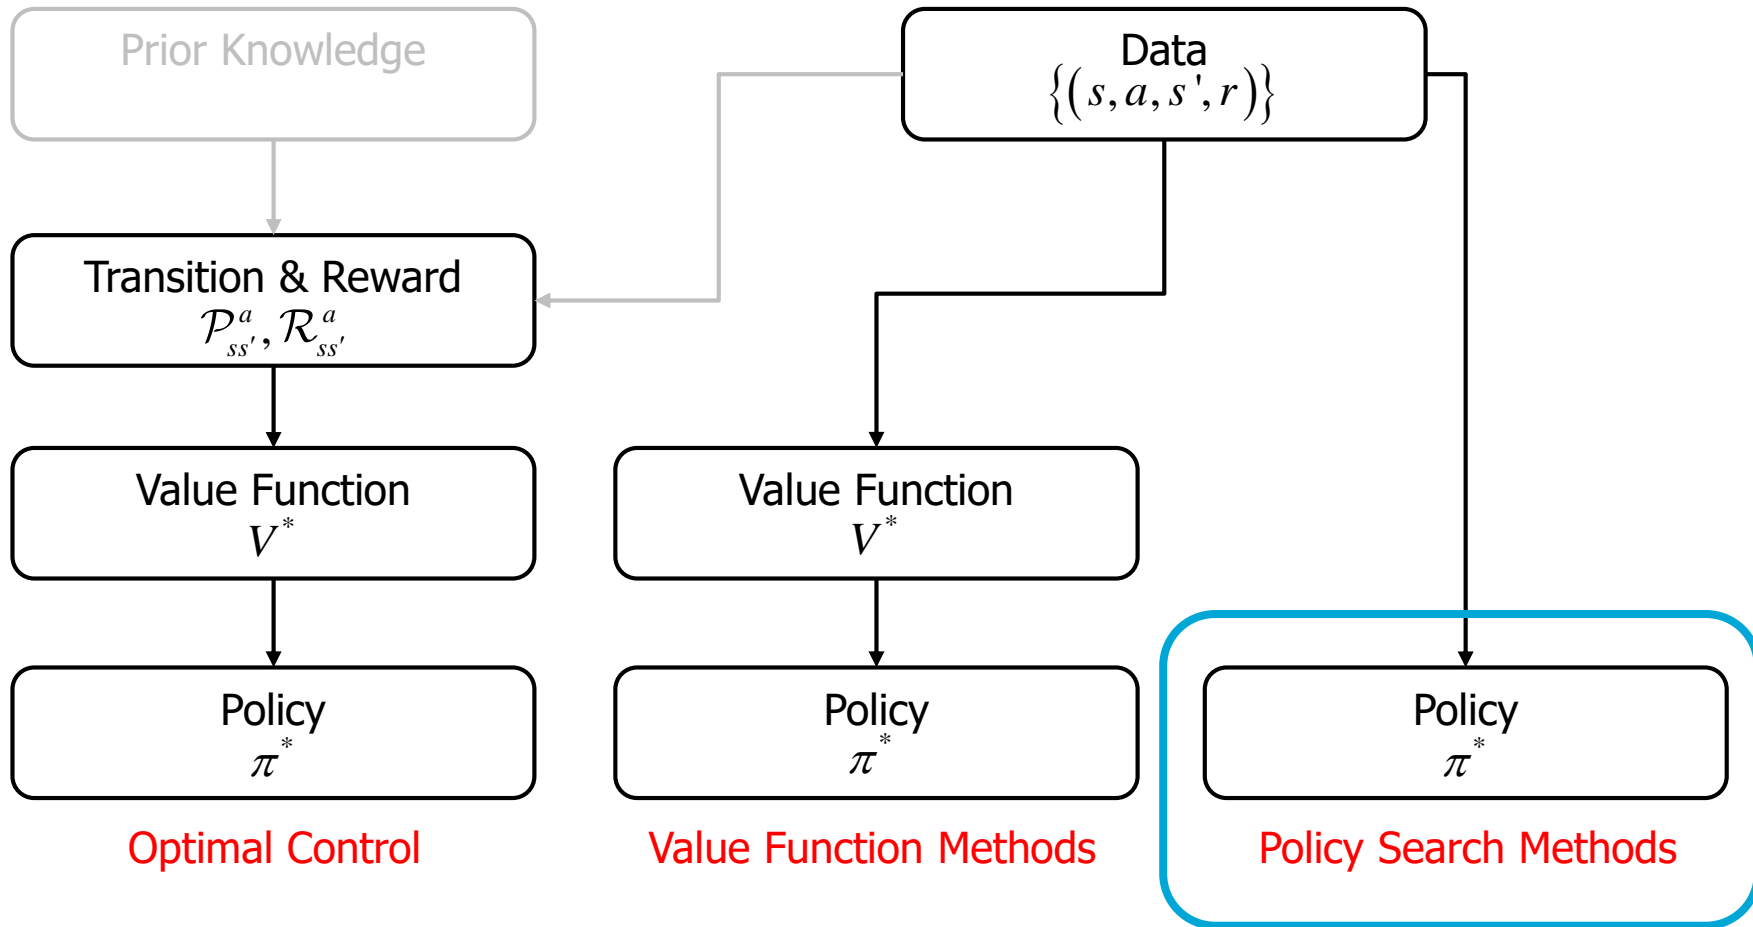
- Function approximation
  - Value function
  - Policy





[https://youtu.be/nM1HTp\\_P3lY](https://youtu.be/nM1HTp_P3lY)

# Types of RL Algorithms



# Policy Search Methods

- **Model-free**
- Local optimum
- Usually parametrized policies
- Examples:
  - Gradient-based
  - Expectation-maximization inspired
  - General optimization



<https://youtu.be/qtqubguikMk>


# Outline

- Introduction
- Notation & basic concepts
- Types of RL algorithms
  - Optimal control
  - Value function methods
  - Policy search
- Summary & outlook

# Summary

## Reinforcement learning

- Inspired by human & animal learning
- Reward as feedback signal
- Model-free, adaptive optimal control



Goal: find  $\pi^*$  maximizing return  $R$

# Challenges of (Robot) RL

- Curse of dimensionality
- Exploration-exploitation trade-off
- Real-world samples
- Model uncertainty
- Goal specification



# Tractability Through:

- Representations
  - State and/or action discretization
  - Value function approximation
  - Pre-structured policies
- Prior knowledge
  - Demonstrations
  - Task Structure
- Models
  - Mental rehearsal